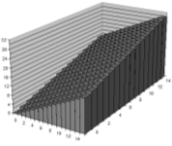


Supplemental Material

Bits, Bytes, and Integer



4190.308 Computer Architecture, Fall 2015

Bits, Bytes, and Integers

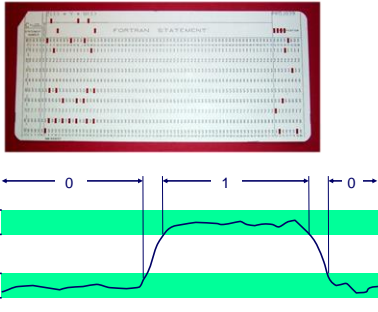
- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- Summary

Acknowledgement: slides based on the cs:app2e material

4190.308 Computer Architecture, Fall 20152

CSE컴퓨터공학부

Binary Representations



4190.308 Computer Architecture, Fall 20153

CSE컴퓨터공학부

Encoding Byte Values


- Byte = 8 bits
- Binary 00000000_2 to 11111111_2
- Decimal: 0_{10} to 255_{10}
- Hexadecimal 00_{16} to FF_{16}
 - Base 16 number representation
 - Use characters '0' to '9' and 'A' to 'F'
 - Write $FA1D37B_{16}$ in C as
 - $0xFA1D37B$
 - $0xfa1d37b$

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

4190.308 Computer Architecture, Fall 20154

CSE컴퓨터공학부

Byte-Oriented Memory Organization



- Programs Refer to Virtual Addresses
 - Conceptually very large array of bytes
 - Actually implemented with hierarchy of different memory types
 - System provides address space private to particular "process"
 - Program being executed
 - Program can clobber its own data, but not that of others
- Compiler + Run-Time System Control Allocation
 - Where different program objects should be stored
 - All allocation within single virtual address space

4190.308 Computer Architecture, Fall 20155

CSE컴퓨터공학부

Machine Words

- Machine Has "Word Size"
- Nominal size of integer-valued data
 - Including addresses
- Most current machines use 32 bits (4 bytes) words
 - Limits addresses to 4GB
 - Becoming too small for memory-intensive applications
- High-end systems use 64 bits (8 bytes) words
 - Potential address space $\approx 1.8 \times 10^{19}$ bytes
 - x86-64 machines support 48-bit addresses: 256 Terabytes
- Machines support multiple data formats
 - Fractions or multiples of word size
 - Always integral number of bytes

4190.308 Computer Architecture, Fall 20156

CSE컴퓨터공학부

Word-Oriented Memory Organization

- Addresses Specify Byte Locations
 - Address of first byte in word
 - Addresses of successive words differ by 4 (32-bit) or 8 (64-bit)

32-bit Words

Addr = 0000

Addr = 0004

Addr = 0008

Addr = 0012

64-bit Words

Addr = 0000

Addr = 0008

Bytes

0000

0001

0002

0003

0004

0005

0006

0007

0008

0009

0010

0011

0012

0013

0014

0015

4190.308 Computer Architecture, Fall 2015

7

CSE

컴퓨터공학부

Data Representations

C Data Type	Typical 32-bit	Intel IA32	x86-64
char	1	1	1
short	2	2	2
int	4	4	4
long	4	4	8
long long	8	8	8
float	4	4	4
double	8	8	8
long double	8	10/12	10/16
pointer	4	4	8

4190.308 Computer Architecture, Fall 2015

8

CSE

컴퓨터공학부

Byte Ordering

- How should bytes within a multi-byte word be ordered in memory?
- Conventions
 - Big Endian: Sun, PPC Mac, Internet
 - Least significant byte has highest address
 - Little Endian: x86
 - Least significant byte has lowest address

4190.308 Computer Architecture, Fall 2015

9

CSE

컴퓨터공학부

Byte Ordering Example

- Big Endian**
 - Least significant byte has highest address
- Little Endian**
 - Least significant byte has lowest address
- Example
 - Variable x has 4-byte representation 0x01234567
 - Address given by &x is 0x100

Big Endian

0x100 0x101 0x102 0x103

01 23 45 67

Little Endian

0x100 0x101 0x102 0x103

67 45 23 01

4190.308 Computer Architecture, Fall 2015

10

CSE

컴퓨터공학부

Reading Byte-Reversed Listings

- Disassembly
 - Text representation of binary machine code
 - Generated by program that reads the machine code
- Example Fragment

Address	Instruction Code	Assembly Rendition
8048365:	5b	pop %ebx
8048366:	81 c3 ab 12 00 00	add \$0x12ab,%ebx
804836c:	83 bb 28 00 00 00	cmpl \$0x0,0x28(%ebx)

0x12ab

0x000012ab

00 00 12 ab

ab 12 00 00

- Deciphering Numbers
 - Value:
 - Pad to 32 bits:
 - Split into bytes:
 - Reverse:

4190.308 Computer Architecture, Fall 2015

11

CSE

컴퓨터공학부

Examining Data Representations

- Code to Print Byte Representation of Data
 - Casting pointer to unsigned char * creates byte array

```
void show_bytes(char *start, int len){  
    int i;  
    for (i = 0; i < len; i++) {  
        printf("%p\t0x%.2x\n", start+i, start[i]);  
    }  
    printf("\n");  
}
```

printf directives:
%p: print pointer
%x: print hexadecimal

4190.308 Computer Architecture, Fall 2015

12

CSE

컴퓨터공학부

show_bytes Execution Example

```
int a = 15213;
printf("int a = 15213;\n");
show_bytes((char *) &a, sizeof(int));
```

Result (Linux):

```
int a = 15213;
0x11ffffcb8 0x6d
0x11ffffcb9 0x3b
0x11ffffcba 0x00
0x11ffffcbb 0x00
```

4190.308 Computer Architecture, Fall 2015

13

CSE 컴퓨터공학부

Representing Integers

Decimal: 15213
Binary: 0011 1011 0110 1101
Hex: 3 B 6 D

int A = 15213;

long int C = 15213;

int B = -15213; (= 0xffffc493)

Two's complement representation (Covered later)

IA32, x86-64 Sun

6D	00
3B	00
00	3B
00	6D

IA32 x86-64 Sun

6D	6D	00
3B	3B	00
00	00	00
00	00	00
00	00	00
00	00	00
00	00	00
00	00	6D

IA32, x86-64 Sun

93	FF
C4	FF
FF	C4
FF	93

4190.308 Computer Architecture, Fall 2015

14

CSE 컴퓨터공학부

Representing Pointers

```
int B = -15213;
int *P = &B;
```

Sun

EF
FF
FB
2C

IA32

D4
F8
FF
BF

x86-64

0C
89
EC
FF
FF
7F
00
00

→ different compilers & machines assign different locations to objects

4190.308 Computer Architecture, Fall 2015

15

CSE 컴퓨터공학부

Representing Strings

```
char S[6] = "18243";
```

Linux/Alpha

31
38
32
34
33
00

Sun

31
38
32
34
33
00

- Strings in C
 - Represented by array of characters
 - Each character encoded in ASCII format
 - Standard 7-bit encoding of character set
 - Character "0" has code 0x30
 - Digit *i* has code 0x30+*i*
 - String should be null-terminated
 - Final character = 0
- Compatibility
 - Byte ordering not an issue

4190.308 Computer Architecture, Fall 2015

16

CSE 컴퓨터공학부

Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- Summary

4190.308 Computer Architecture, Fall 2015

17

CSE 컴퓨터공학부

Boolean Algebra

- Developed by George Boole in 19th Century
- Algebraic representation of logic
 - Encode "True" as 1 and "False" as 0

And

A & B = 1 when both A=1 and B=1

&	0	1
0	0	0
1	0	1

Not

~A = 1 when A=0

~	0	1
0	1	
1	0	

Or

A | B = 1 when either A=1 or B=1

	0	1
0	0	1
1	1	1

Exclusive-Or (Xor)

A ^ B = 1 when either A=1 or B=1, but not both

^	0	1
0	0	1
1	1	0

4190.308 Computer Architecture, Fall 2015

18

CSE 컴퓨터공학부

Application of Boolean Algebra

■ Applied to Digital Systems by Claude Shannon

● 1937 MIT Master's Thesis

● Reason about networks of relay switches

▶ Encode closed switch as 1, open switch as 0

A

~A

~B

B

A&~B

~A&B

Connection when

$A \& \sim B \mid \sim A \& B$

$= A \wedge B$

4190.308 Computer Architecture, Fall 2015

19

CSE

컴퓨터공학부

Department of Computer Science & Engineering

General Boolean Algebras

■ Operate on Bit Vectors

● Operations applied bitwise

01101001

$\&$ 01010101

01000001

01101001

\mid 01010101

01111101

01101001

\wedge 01010101

00111100

01010101

\sim 01010101

10101010

■ All of the Properties of Boolean Algebra Apply

4190.308 Computer Architecture, Fall 2015

20

CSE

컴퓨터공학부

Department of Computer Science & Engineering

Representing & Manipulating Sets

■ Representation

● Width w bit vector represents subsets of $\{0, \dots, w-1\}$

● $a_j = 1$ if $j \in A$

▶ 01101001 $\{0, 3, 5, 6\}$

▶ 76543210

▶ 01010101 $\{0, 2, 4, 6\}$

▶ 76543210

■ Operations

&

Intersection

01000001

$\{0, 6\}$

|

Union

01111101

$\{0, 2, 3, 4, 5, 6\}$

^

Symmetric difference

00111100

$\{2, 3, 4, 5\}$

~

Complement

10101010

$\{1, 3, 5, 7\}$

4190.308 Computer Architecture, Fall 2015

21

CSE

컴퓨터공학부

Department of Computer Science & Engineering

Bit-Level Operations in C

■ Operations $\&$, \mid , \sim , \wedge available in C

● Apply to any "integral" data type

▶ long, int, short, char, unsigned

● View arguments as bit vectors

● Arguments applied bit-wise

■ Examples (Char data type)

● $\sim 0x41 \rightarrow 0xBE$

▶ $\sim 01000001_2 \rightarrow 10111110_2$

● $\sim 0x00 \rightarrow 0xFF$

▶ $\sim 00000000_2 \rightarrow 11111111_2$

● $0x69 \& 0x55 \rightarrow 0x41$

▶ $01101001_2 \& 01010101_2 \rightarrow 01000001_2$

● $0x69 \mid 0x55 \rightarrow 0x7D$

▶ $01101001_2 \mid 01010101_2 \rightarrow 01111101_2$

4190.308 Computer Architecture, Fall 2015

22

CSE

컴퓨터공학부

Department of Computer Science & Engineering

Contrast: Logic Operations in C

■ Contrast to Logical Operators

● $\&$, \mid , $!$

▶ View 0 as "False"

▶ Anything nonzero as "True"

▶ Always return 0 or 1

▶ Early termination

■ Examples (char data type)

● $!0x41 \rightarrow 0x00$

● $!0x00 \rightarrow 0x01$

● $!!0x41 \rightarrow 0x01$

● $0x69 \&\& 0x55 \rightarrow 0x01$

● $0x69 \mid\mid 0x55 \rightarrow 0x01$

● $p \&\& *p$ (avoids null pointer access)

4190.308 Computer Architecture, Fall 2015

23

CSE

컴퓨터공학부

Department of Computer Science & Engineering

Shift Operations

■ Left Shift: $x \ll y$

● Shift bit-vector x left y positions

– Throw away extra bits on left

▶ Fill with 0's on right

■ Right Shift: $x \gg y$

● Shift bit-vector x right y positions

– Throw away extra bits on right

● Logical shift

▶ Fill with 0's on left

● Arithmetic shift

▶ Replicate most significant bit on right

■ Undefined Behavior

● Shift amount < 0 or \geq word size

Argument x

01100010

$\ll 3$

00010000

Log. $\gg 2$

00011000

Arith. $\gg 2$

00011000

Argument x

10100010

$\ll 3$

00010000

Log. $\gg 2$

00101000

Arith. $\gg 2$

11101000

CSE

컴퓨터공학부

Department of Computer Science & Engineering

4

Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- Summary

Encoding Integers

Unsigned

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

Two's Complement

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

short int x = 15213;
short int y = -15213;

Sign Bit

- C short 2 bytes long

	Decimal	Hex	Binary
x	15213	3B 6D	00111011 01101101
y	-15213	C4 93	11000100 10010011
- Sign Bit
 - For 2's complement, most significant bit indicates sign
 - 0 for nonnegative
 - 1 for negative

Encoding Example (Cont.)

x = 15213: 00111011 01101101
y = -15213: 11000100 10010011

Weight	15213	-15213
1	1	1
2	0	1
4	1	0
8	1	0
16	0	1
32	1	0
64	1	0
128	0	1
256	1	0
512	1	0
1024	0	1
2048	1	0
4096	1	0
8192	1	0
16384	0	1
-32768	0	1
Sum	15213	-15213

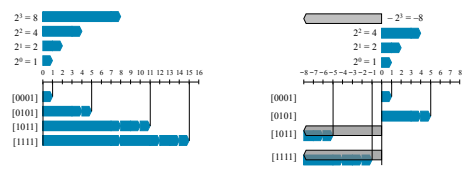
Encoding Example Visualized

Unsigned

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

Two's Complement

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$



Numeric Ranges

- Unsigned Values
 - UMin = 0
000...0
 - UMax = 2^w - 1
111...1
- Two's Complement Values
 - TMin = -2^{w-1}
100...0
 - TMax = 2^{w-1} - 1
011...1
- Other Values
 - Minus 1
111...1

Values for W = 16

	Decimal	Hex	Binary
UMax	65535	FF FF	11111111 11111111
TMax	32767	7F FF	01111111 11111111
TMin	-32768	80 00	10000000 00000000
-1	-1	FF FF	11111111 11111111
0	0	00 00	00000000 00000000

Values for Different Word Sizes

	W			
	8	16	32	64
UMax	255	65,535	4,294,967,295	18,446,744,073,709,551,615
TMax	127	32,767	2,147,483,647	9,223,372,036,854,775,807
TMin	-128	-32,768	-2,147,483,648	-9,223,372,036,854,775,808

- Observations
 - |TMin| = TMax + 1
 - Asymmetric range
 - UMax = 2 * TMax + 1
- C Programming
 - #include <limits.h>
 - Declares constants, e.g.,
 - ULONG_MAX
 - LONG_MAX
 - LONG_MIN
 - Values platform specific

Unsigned & Signed Numeric Values

X	B2U(X)	B2T(X)
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	8	-8
1001	9	-7
1010	10	-6
1011	11	-5
1100	12	-4
1101	13	-3
1110	14	-2
1111	15	-1

- Equivalence
 - Same encodings for nonnegative values
- Uniqueness
 - Every bit pattern represents unique integer value
 - Each representable integer has unique bit encoding
- ⇒ Can Invert Mappings
 - $U2B(x) = B2U^{-1}(x)$
 - Bit pattern for unsigned integer
 - $T2B(x) = B2T^{-1}(x)$
 - Bit pattern for two's comp integer

4190.308 Computer Architecture, Fall 2015

31

CSE 컴퓨터공학부

An (Easier) Way To Look At It

Unsigned
$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

Two's Complement
$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

4190.308 Computer Architecture, Fall 2015

32

CSE 컴퓨터공학부

Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - **Conversion, casting**
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- Summary

4190.308 Computer Architecture, Fall 2015

33

CSE 컴퓨터공학부

Mapping Between Signed & Unsigned

Two's Complement
$$x \xrightarrow{\begin{matrix} T2U \\ \text{B2U} \end{matrix}} u_x$$

Unsigned
$$u_x \xrightarrow{\begin{matrix} U2T \\ \text{B2T} \end{matrix}} x$$

Maintain Same Bit Pattern

- Mappings between unsigned and two's complement numbers: keep bit representations and reinterpret

4190.308 Computer Architecture, Fall 2015

34

CSE 컴퓨터공학부

Mapping Signed ↔ Unsigned

Bits	Signed	Unsigned
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	-8	8
1001	-7	9
1010	-6	10
1011	-5	11
1100	-4	12
1101	-3	13
1110	-2	14
1111	-1	15

Arrows indicate T2U (Signed to Unsigned) and U2T (Unsigned to Signed) mappings.

4190.308 Computer Architecture, Fall 2015

35

CSE 컴퓨터공학부

Mapping Signed ↔ Unsigned

Bits	Signed	Unsigned
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	-8	8
1001	-7	9
1010	-6	10
1011	-5	11
1100	-4	12
1101	-3	13
1110	-2	14
1111	-1	15

Arrows indicate the mapping between Signed and Unsigned values, with a note for the wrap-around from -8 to 8 (+/- 16).

4190.308 Computer Architecture, Fall 2015

36

CSE 컴퓨터공학부

Relation between Signed & Unsigned

Two's Complement

Unsigned

x → ux

T2U
T2B
B2U

Maintain Same Bit Pattern

$w-1$ 0

ux x

Large negative weight becomes Large positive weight

$ux = \begin{cases} x & x \geq 0 \\ x + 2^w & x < 0 \end{cases}$

4190.308 Computer Architecture, Fall 201537CSE 컴퓨터공학부

Conversion Visualized

2's Comp. → Unsigned

- Ordering Inversion
- Negative → Big Positive

2's Complement Range

Unsigned Range

$TMax$ $UMax$

0 $UMax - 1$

-1 $TMax + 1$

-2 $TMax$

$TMin$ 0

4190.308 Computer Architecture, Fall 201538CSE 컴퓨터공학부

Conversion Visualized (2)

2's Comp. → Unsigned

- Ordering Inversion
- Negative → Big Positive

Unsigned → 2's Comp.

- Ordering Inversion
- Big Positive → Negative

Conversion

- Number $\leq TMax$: Unsigned == 2's Comp
- Otherwise: Unsigned = 2's Comp + 2^w

$UMax$

0

-1

$TMin$ $TMax$

± 16

$==$

4190.308 Computer Architecture, Fall 201539CSE 컴퓨터공학부

Signed vs. Unsigned in C

Constants

- By default are considered to be signed integers
- Unsigned if have "U" as suffix
 - 0U, 4294967259U

Casting

- Explicit casting between signed & unsigned same as U2T and T2U
 - int tx, ty;
 - unsigned ux, uy;
 - tx = (int) ux;
 - uy = (unsigned) ty;
- Implicit casting also occurs via assignments and procedure calls
 - tx = ux;
 - uy = ty;

4190.308 Computer Architecture, Fall 201540CSE 컴퓨터공학부

Casting Surprises

Expression Evaluation

- If there is a mix of unsigned and signed in single expression, **signed values implicitly cast to unsigned**
- Including comparison operations $<$, $>$, $==$, $<=$, $>=$
- Examples for $W = 32$: $TMin = -2,147,483,648$, $TMax = 2,147,483,647$

Constant ₁	Constant ₂	Relation	Evaluation
0	0U	==	unsigned
-1	0	<	signed
-1	0U	>	unsigned
2147483647	-2147483647-1	>	signed
2147483647U	-2147483647-1	<	unsigned
-1	-2	>	signed
(unsigned)-1	-2	>	unsigned
2147483647	2147483648U	<	unsigned
2147483647	(int) 2147483648U	>	signed

4190.308 Computer Architecture, Fall 201541CSE 컴퓨터공학부

Code Security Example

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}

/* memcpy definition */
void *memcpy(void *dest, const void *src, size_t n);
```

- Similar to code found in FreeBSD's implementation of getpeername
- There are legions of smart people trying to find vulnerabilities in programs

4190.308 Computer Architecture, Fall 201542CSE 컴퓨터공학부

Typical Usage

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}

/* memcpy definition */
void *memcpy(void *dest, const void *src, size_t n);

#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, MSIZE);
    printf("%s\n", mybuf);
}
```

Malicious Usage

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}

/* memcpy definition */
void *memcpy(void *dest, const void *src, size_t n);

#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, -MSIZE);
    ...
}
```

Understanding What Can Go Wrong

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}

/* memcpy definition */
void *memcpy(void *dest, const void *src, size_t n);

#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, -MSIZE);
    ...
}
```

Summary

Casting Signed ↔ Unsigned: Basic Rules


- Bit pattern is maintained
- But reinterpreted
- Can have unexpected effects: adding or subtracting 2^w
- Expression containing signed and unsigned int
 - int is cast to unsigned!!

Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- Summary

Sign Extension

- Task:
 - Given w-bit signed integer x
 - Convert it to w+k-bit integer with same value
- Rule:
 - Make k copies of sign bit:
 - $X' = x_{w-1}, \dots, x_{w-1}, x_{w-1}, x_{w-2}, \dots, x_0$



Sign Extension Example

```
short int x = 15213;
int ix = (int) x;
short int y = -15213;
int iy = (int) y;
```

	Decimal	Hex	Binary
x	15213	3B 6D	00111011 01101101
ix	15213	00 00 3B 6D	00000000 00000000 00111011 01101101
y	-15213	C4 93	11000100 10010011
iy	-15213	FF FF C4 93	11111111 11111111 11000100 10010011

- Converting from smaller to larger integer data type
- C automatically performs sign extension

Summary: Expanding, Truncating: Basic Rules

- Expanding (e.g., short int to int)
 - Unsigned: zeros added
 - Signed: sign extension
 - Both yield expected result
- Truncating (e.g., unsigned to unsigned short)
 - Unsigned/signed: bits are truncated
 - Result reinterpreted
 - Unsigned: mod operation
 - Signed: similar to mod
 - For small numbers yields expected behavior

Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- Summary

Negation: Complement & Increment

- Claim: Following Holds for 2's Complement
$$\sim x + 1 == -x$$
- Complement
 - Observation: $\sim x + x == 1111...111 == -1$

$$\begin{array}{r} x \quad 10011101 \\ + \sim x \quad 01100010 \\ \hline -1 \quad 11111111 \end{array}$$

Complement & Increment Examples

x = 15213

	Decimal	Hex	Binary
x	15213	3B 6D	00111011 01101101
~x	-15214	C4 92	11000100 10010010
~x+1	-15213	C4 93	11000100 10010011
y	-15213	C4 93	11000100 10010011

x = 0

	Decimal	Hex	Binary
0	0	00 00	00000000 00000000
~0	-1	FF FF	11111111 11111111
~0+1	0	00 00	00000000 00000000

Unsigned Addition

Operands: w bits

True Sum: w+1 bits

Discard Carry: w bits


$$\begin{array}{r} u \quad \boxed{} \boxed{} \dots \boxed{} \boxed{} \\ + v \quad \boxed{} \boxed{} \dots \boxed{} \boxed{} \\ \hline u+v \quad \boxed{} \boxed{} \dots \boxed{} \boxed{} \\ \hline \text{UAdd}_w(u, v) \quad \boxed{} \boxed{} \dots \boxed{} \boxed{} \end{array}$$

- Standard Addition Function
 - Ignores carry output
- Implements Modular Arithmetic
$$s = \text{UAdd}_w(u, v) = u + v \bmod 2^w$$

$$\text{UAdd}_w(u, v) = \begin{cases} u + v & u + v < 2^w \\ u + v - 2^w & u + v \geq 2^w \end{cases}$$



Visualizing (Mathematical) Integer Addition

- Integer Addition
 - 4-bit integers u, v
 - Compute true sum $Add_4(u, v)$
 - Values increase linearly with u and v
 - Forms planar surface



Visualizing Unsigned Addition

- Wraps Around
 - If true sum $\geq 2^w$
 - At most once



Mathematical Properties

- Modular Addition Forms an *Abelian Group*
 - Closed** under addition
 - $0 \leq UAdd_w(u, v) \leq 2^w - 1$
 - Commutative**
 - $UAdd_w(u, v) = UAdd_w(v, u)$
 - Associative**
 - $UAdd_w(t, UAdd_w(u, v)) = UAdd_w(UAdd_w(t, u), v)$
 - 0** is additive identity
 - $UAdd_w(u, 0) = u$
 - Every element has additive **inverse**
 - Let $UComp_w(u) = 2^w - u$
 - $UAdd_w(u, UComp_w(u)) = 0$

Two's Complement Addition




- TAdd and UAdd have Identical Bit-Level Behavior
 - Signed vs. unsigned addition in C:

```
int s, t, u, v;
s = (int) ((unsigned) u + (unsigned) v);
t = u + v;
```
 - Will give `s == t`


TAdd Overflow

- Functionality
 - True sum requires $w+1$ bits
 - Drop off MSB
 - Treat remaining bits as 2's comp. integer



Visualizing 2's Complement Addition

- Values
 - 4-bit two's comp.
 - Range from -8 to +7
- Wraps Around
 - If $sum \geq 2^{w-1}$
 - Becomes negative
 - At most once
 - If $sum < -2^{w-1}$
 - Becomes positive
 - At most once



Characterizing TAdd

■ Functionality

- True sum requires $w+1$ bits
- Drop off MSB
- Treat remaining bits as 2's comp. integer

$TAdd(u, v)$

Positive Overflow

Negative Overflow

> 0

< 0

v

$u > 0$

$TAdd_w(u, v) = \begin{cases} u + v + 2^w & u + v < TMin_w \text{ (NegOver)} \\ u + v & TMin_w \leq u + v \leq TMax_w \\ u + v - 2^w & TMax_w < u + v \text{ (PosOver)} \end{cases}$

Mathematical Properties of TAdd

■ Isomorphic Group to unsigneds with UAdd

- $TAdd_w(u, v) = U2T(UAdd_w(T2U(u), T2U(v)))$
 - Since both have identical bit patterns

■ Two's Complement Under TAdd Forms a Group

- Closed, Commutative, Associative, 0 is additive identity
- Every element has additive inverse

$TComp_w(u) = \begin{cases} -u & u \neq TMin_w \\ TMin_w & u = TMin_w \end{cases}$

Multiplication

■ Computing Exact Product of w -bit numbers x, y

- Either signed or unsigned

■ Ranges

- Unsigned: $0 \leq x * y \leq (2^w - 1)^2 = 2^{2w} - 2^{w+1} + 1$
 - Up to $2w$ bits
- Two's complement min: $x * y \geq (-2^{w-1}) * (2^{w-1} - 1) = -2^{2w-2} + 2^{w-1}$
 - Up to $2w-1$ bits
- Two's complement max: $x * y \leq (2^{w-1} - 1)^2 = 2^{2w-2} - 2^{w-1} + 1$
 - Up to $2w$ bits, but only for $(TMin_w)^2$

■ Maintaining Exact Results

- Would need to keep expanding word size with each product computed
- Done in software by "arbitrary precision" arithmetic packages

Unsigned Multiplication in C

Operands: w bits

u

v

$*$

True Product: 2^w bits

$u \cdot v$

Discard w bits: w bits

$UMult_w(u, v)$

■ Standard Multiplication Function

- Ignores high order w bits

■ Implements Modular Arithmetic

$UMult_w(u, v) = u \cdot v \bmod 2^w$

Signed Multiplication in C

Operands: w bits

u

v

$*$

True Product: 2^w bits

$u \cdot v$

Discard w bits: w bits

$TMult_w(u, v)$

■ Standard Multiplication Function

- Ignores high order w bits
- Some of which are different for signed vs. unsigned multiplication
- Lower bits are the same

Power-of-2 Multiply with Shift

■ Operation

- $u \ll k$ gives $u * 2^k$
- Both signed and unsigned

Operands: w bits

u

2^k

$*$

True Product: $w+k$ bits

$u \cdot 2^k$

Discard k bits: w bits

$UMult_w(u, 2^k)$

$TMult_w(u, 2^k)$

■ Examples

- $u \ll 3 \quad == \quad u * 8$
- $u \ll 5 - u \ll 3 \quad == \quad u * 24$
- Most machines shift and add faster than multiply
 - Compiler generates this code automatically

Compiled Multiplication Code

C Function

```
int mul12(int x)
{
    return x*12;
}
```

Compiled Arithmetic Operations

```
leal (%eax,%eax,2), %eax
sall $2, %eax
```


Explanation

```
t <- x+x*2
return t << 2;
```

- C compiler automatically generates shift/add code when multiplying by a constant

Unsigned Power-of-2 Divide with Shift

- Quotient of Unsigned by Power of 2
 - $u \gg k$ gives $\lfloor u / 2^k \rfloor$
 - Uses logical shift



	Division	Computed	Hex	Binary
x	15213	15213	3B 6D	00111011 01101101
x >> 1	7606.5	7606	1D B6	00011101 10110110
x >> 4	950.8125	950	03 B6	00000011 10110110
x >> 8	59.4257813	59	00 3B	00000000 00111011

Compiled Unsigned Division Code

C Function

```
unsigned udiv8(unsigned x)
{
    return x/8;
}
```

Compiled Arithmetic Operations

```
shrl $3, %eax
```


Explanation

```
# Logical shift
return x >> 3;
```

- Uses logical shift for unsigned
- For Java Users
 - Logical shift written as >>>

Signed Power-of-2 Divide with Shift

- Quotient of Signed by Power of 2
 - $x \ggg k$ gives $\lfloor x / 2^k \rfloor$
 - Uses arithmetic shift
 - Rounds wrong direction when $u < 0$




	Division	Computed	Hex	Binary
y	-15213	-15213	C4 93	11000100 10010011
y >> 1	-7606.5	-7607	E2 49	11100010 01001001
y >> 4	-950.8125	-951	FC 49	11111100 01001001
y >> 8	-59.4257813	-60	FF C4	11111111 11000100

Correct Power-of-2 Divide

- Quotient of Negative Number by Power of 2
 - Want $\lceil x / 2^k \rceil$ (Round Toward 0)
 - Compute as $\lfloor (x+2^k-1) / 2^k \rfloor$
 - In C: $(x + (1 < k) - 1) \ggg k$
 - Biases dividend toward 0


Case 1: No rounding



Biasing has no effect

Correct Power-of-2 Divide (Cont.)

Case 2: Rounding



Biasing adds 1 to final result

Compiled Signed Division Code

C Function

```
int idiv8(int x)
{
    return x/8;
}
```

Compiled Arithmetic Operations

```
testl %eax, %eax
js     L4
L3:    sarl $3, %eax
ret
L4:    addl $7, %eax
jmp    L3
```


Explanation

```
if x < 0
    x += 7;
# Arithmetic shift
return x >> 3;
```

- Uses arithmetic shift for int
- For Java Users
 - Arith. shift written as >>

4190.308 Computer Architecture, Fall 2015

73


 CSE 컴퓨터공학부
Department of Computer Science & Engineering

Arithmetic: Basic Rules

- Addition:
 - Unsigned/signed: Normal addition followed by truncate, same operation on bit level
 - Unsigned: addition mod 2^w
 - ▶ Mathematical addition + possible subtraction of 2^w
 - Signed: modified addition mod 2^w (result in proper range)
 - ▶ Mathematical addition + possible addition or subtraction of 2^w
- Multiplication:
 - Unsigned/signed: Normal multiplication followed by truncate, same operation on bit level
 - Unsigned: multiplication mod 2^w
 - Signed: modified multiplication mod 2^w (result in proper range)

4190.308 Computer Architecture, Fall 2015

74


 CSE 컴퓨터공학부
Department of Computer Science & Engineering

Arithmetic: Basic Rules

- Unsigned ints, 2's complement ints are isomorphic rings: isomorphism = casting
- Left shift
 - Unsigned/signed: multiplication by 2^k
 - Always logical shift
- Right shift
 - Unsigned: logical shift, div (division + round to zero) by 2^k
 - Signed: arithmetic shift
 - ▶ Positive numbers: div (division + round to zero) by 2^k
 - ▶ Negative numbers: div (division + round away from zero) by 2^k
 - Use biasing to fix

4190.308 Computer Architecture, Fall 2015

75


 CSE 컴퓨터공학부
Department of Computer Science & Engineering

Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- Summary

4190.308 Computer Architecture, Fall 2015

76


 CSE 컴퓨터공학부
Department of Computer Science & Engineering

Properties of Unsigned Arithmetic

- Unsigned Multiplication with Addition Forms Commutative Ring
 - Addition is commutative group
 - Closed under multiplication
$$0 \leq \text{UMult}_w(u, v) \leq 2^w - 1$$
 - Multiplication Commutative
$$\text{UMult}_w(u, v) = \text{UMult}_w(v, u)$$
 - Multiplication is Associative
$$\text{UMult}_w(t, \text{UMult}_w(u, v)) = \text{UMult}_w(\text{UMult}_w(t, u), v)$$
 - 1 is multiplicative identity
$$\text{UMult}_w(u, 1) = u$$
 - Multiplication distributes over addition
$$\text{UMult}_w(t, \text{UAdd}_w(u, v)) = \text{UAdd}_w(\text{UMult}_w(t, u), \text{UMult}_w(t, v))$$

4190.308 Computer Architecture, Fall 2015

77


 CSE 컴퓨터공학부
Department of Computer Science & Engineering

Properties of Two's Comp. Arithmetic

- Isomorphic Algebras
 - Unsigned multiplication and addition
 - ▶ Truncating to w bits
 - Two's complement multiplication and addition
 - ▶ Truncating to w bits
- Both Form Rings
 - Isomorphic to ring of integers mod 2^w
- Comparison to (Mathematical) Integer Arithmetic
 - Both are rings
 - Integers obey ordering properties, e.g.,
$$u > 0 \quad \Rightarrow \quad u + v > v$$
$$u > 0, v > 0 \quad \Rightarrow \quad u \cdot v > 0$$
 - These properties are not obeyed by two's comp. arithmetic
$$TMax + 1 \quad == \quad TMin$$
$$15213 * 30426 \quad == \quad -10030 \text{ (16-bit words)}$$

4190.308 Computer Architecture, Fall 2015

78

 CSE 컴퓨터공학부
Department of Computer Science & Engineering

Why Should I Use Unsigned?

- *Don't* Use Just Because Number Nonnegative
 - Easy to make mistakes

```
unsigned i;
for (i = cnt-2; i >= 0; i--)
    a[i] += a[i+1];
```
 - Can be very subtle

```
#define DELTA sizeof(int)
int i;
for (i = CNT; i-DELTA >= 0; i-= DELTA)
    . . .
```
- Do Use When Performing Modular Arithmetic
 - Multiprecision arithmetic
- Do Use When Using Bits to Represent Sets
 - Logical right shift, no sign extension

4190.308 Computer Architecture, Fall 2015

79

 CSE 컴퓨터공학부
Department of Computer Science & Engineering