

The HW/SW Interface

The x86 ISA:
Machine-Level Programming Basics

```
0x08048394 <+0>: push    %ebp
0x08048395 <+1>: mov     %esp, %ebp
0x08048397 <+3>: mov     0xc(%ebp), %eax
0x0804839a <+6>: add     0x8(%ebp), %eax
0x0804839d <+9>: pop     %ebp
0x0804839e <+10>: ret
0x08048394 <sum>: 0x55    0x89    0xe5 ...
0x0804839c <sum+8>: 0x08    0x5d    0xc3
```

4190.308 Computer Architecture, Fall 2015

Instruction Set Architecture (ISA)

ISA classification

- RISC vs CISC
- general-purpose
 - two or three operands?
 - # of memory operands?
 - classification
 - load-store
 - register-memory
 - memory-memory
- stack
- accumulator

Diagram illustrating ISA classification. It shows a stack architecture (a) where data is pushed/popped from a stack. An accumulator architecture (b) where data is stored in a single register (the accumulator). Register-memory architectures (c) where data is moved between registers and memory. Register-register architectures (d) where data is moved between registers. The diagram also shows a processor with ALUs and registers, and a memory hierarchy.

4190.308 Computer Architecture, Fall 2015

Instruction Set Architecture (ISA)

The ISA defines the **programmer-visible state** of a processor architecture

Diagram illustrating the programmer-visible state of a processor architecture. It shows the state of the processor, including the program counter (PC), instruction register (IR), and the state of the registers (RAX, RBX, RCX, RDX, RSI, RDI, RBP, RSP, R8, R9, R10, R11, R12, R13, R14, R15). The diagram also shows the state of the memory (RAM) and the state of the cache (L1, L2, L3).

- the ISA tells us what we can expect the processor to do when we execute instructions (but not *how* the processor does it)
- memory addressing
 - interpretation of addresses
 - bytes, half words, words, double words, ...?
 - byte ordering in memory
 - little-endian vs. big-endian
 - addressing modes
 - how to access objects in memory

4190.308 Computer Architecture, Fall 2015

Instruction Set Architecture (ISA)

The ISA defines the **programmer-visible state** of a processor architecture

- type and size of operands
 - specifying types
 - byte, half word, word, float, double...
- operations
 - arithmetic and logical
 - data transfer
 - control
 - system
 - floating point
 - string
 - multimedia

Rank	x86 operation	%
1	load	22
2	conditional branch	20
3	compare	16
4	store	12
5	add	8
6	and	6
7	sub	5
8	move reg-reg	4
9	call	1
10	return	1
	total	96

4190.308 Computer Architecture, Fall 2015

Instruction Set Architecture (ISA)

The ISA defines the **programmer-visible state** of a processor architecture

- instruction encoding
 - defines binary representation for each operation and operand
 - fixed size vs. variable size

Operation and no. of operands	Address specifier 1	Address field 1	...	Address specifier n	Address field n
(a) Variable (e.g., Intel 80x86, VAX)					

Operation	Address field 1	Address field 2	Address field 3
(b) Fixed (e.g., Alpha, ARM, MIPS, PowerPC, SPARC, SuperH)			

4190.308 Computer Architecture, Fall 2015

x86 Machine-Level Programming Basics

History of Intel processors and architectures

- C, assembly, machine code
- Assembly Basics: Registers, operands, move, push/pop
- Introduction to x86-64

Acknowledgement: slides based on the cs.app2e material

4190.308 Computer Architecture, Fall 2015

CSE 컴퓨터공학부

Department of Computer Science & Engineering

1

Intel x86 Processors

■ Dominate laptop/desktop/server market

■ Evolutionary design

- Backwards compatible up until 8086, introduced in 1978
- Added more features as time goes on

■ Complex instruction set computer (CISC)

- Many different instructions with many different formats
 - But, only small subset encountered with Linux programs
- Hard to match performance of Reduced Instruction Set Computers (RISC)
- But, Intel has done just that!
 - In terms of speed. Less so for low power.

4190.308 Computer Architecture, Fall 2015

7

CSE

컴퓨터공학부

Department of Computer Science & Engineering

Intel x86 Evolution: Milestones

Name	Date	Transistors	MHz
■ 8086	1978	29K	5-10
• First 16-bit processor. Basis for IBM PC & DOS			
• 1MB address space			
■ 386	1985	275K	16-33
• First 32 bit processor , referred to as IA32			
• Added "flat addressing"			
• Capable of running Unix			
• 32-bit Linux/gcc uses no instructions introduced in later models			
■ Pentium 4F	2004	125M	2800-3800
• First 64-bit processor, referred to as x86-64			
■ Core i7	2008	731M	2667-3333
• multiple cores			

4190.308 Computer Architecture, Fall 2015

8

CSE

컴퓨터공학부

Department of Computer Science & Engineering

Intel x86 Processors: Overview

Architectures

X86-16

X86-32/IA32

MMX

SSE

SSE2

SSE3

X86-64 / EM64t

SSE4

Processors

8086

286

386

486

Pentium

Pentium MMX

Pentium III

Pentium 4

Pentium 4E

Pentium 4F

Core 2 Duo

Core i7

time

IA: often redefined as latest Intel architecture

4190.308 Computer Architecture, Fall 2015

9

CSE

컴퓨터공학부

Department of Computer Science & Engineering

Intel x86 Processors

■ Machine Evolution

• 386	1985	0.3M
• 486	1989	1.9M
• Pentium	1993	3.1M
• Pentium/MMX	1997	4.5M
• PentiumPro	1995	6.5M
• Pentium III	1999	8.2M
• Pentium 4	2001	42M
• Core 2 Duo	2006	291M
• Core i7 (4 cores)	2008	731M
• Core i7 (6 cores)	2011	2'270M
• Xeon E7 v2 (15 c)	2014	4'310M

■ Added Features

- Instructions to support multimedia operations
 - Parallel operations on 1, 2, and 4-byte data, both integer & FP
- Instructions to enable more efficient conditional operations

Integrated Memory Controller ~ 3 Ch DDR3

Core 0 Core 1 Core 2 Core 3

Shared L3 Cache

Core i7 (45nm)

4190.308 Computer Architecture, Fall 2015

10

CSE

컴퓨터공학부

Department of Computer Science & Engineering

New Species: ia64, then IPF, then Itanium,...

Name	Date	Transistors
■ Itanium	2001	10M
• First shot at 64-bit architecture: first called IA64		
• Radically new instruction set designed for high performance		
• Can run existing IA32 programs <ul style="list-style-type: none">▸ On-board "x86 engine"		
• Joint project with Hewlett-Packard		
■ Itanium 2	2002	221M
• Big performance boost		
■ Itanium 2 Dual-Core	2006	1.7B
■ Itanium has not taken off in marketplace <ul style="list-style-type: none">• very fast (esp. FP), very hot, and very expensive• Lack of backward compatibility, no good compiler support, Pentium 4 got too good, overtaken by 64-bit x86 designs		

4190.308 Computer Architecture, Fall 2015

11

CSE

컴퓨터공학부

Department of Computer Science & Engineering

x86 Clones: Advanced Micro Devices (AMD)

■ Historically

- AMD has followed just behind Intel
- A little bit slower, a lot cheaper

■ Then

- Recruited top circuit designers from Digital Equipment Corp. and other downward trending companies
- Built Opteron: tough competitor to Pentium 4
- Developed x86-64, their own extension to 64 bits

■ Recently

- Intel much quicker with multi-core design
- Intel currently far ahead in performance
- Intel em64t backwards compatible to x86-64

4190.308 Computer Architecture, Fall 2015

12

CSE

컴퓨터공학부

Department of Computer Science & Engineering

Intel's 64-Bit

- Intel Attempted Radical Shift from IA32 to IA64
 - Totally different architecture (Itanium)
 - Executes IA32 code only as legacy
 - Application performance disappointing
- AMD Stepped in with Evolutionary Solution
 - x86-64 (now called "AMD64")
- Intel Felt Obligated to Focus on IA64
 - Hard to admit mistake or that AMD is better
- 2004: Intel Announces EM64T extension to IA32
 - Extended Memory 64-bit Technology
 - Almost identical to x86-64!

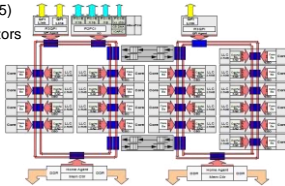
4190.308 Computer Architecture, Fall 2015

13

CSE컴퓨터공학부

Intel Haswell

- Current leader (soon to be replace)
 - Xeon E7v3 (E7-8890v3) (May 2015)
 - 18-core chip: 5.5 billion transistors
 - 2.5/3.3GHz
 - 2 threads per core
 - no GPU on chip
 - virtualization support



source: The Platform

- 18 cores, 36 threads
- per core: 64KB L1, 256KB L2 cache
- 45MB shared L3 cache
- 165W TPD

4190.308 Computer Architecture, Fall 2015

source: The Platform

14

CSE컴퓨터공학부

Our Coverage

- IA32
 - The traditional x86
- x86-64/EM64T
 - The emerging standard
- Presentation
 - ia32 in sections 3.1—3.7 of the textbook
 - x86-64 in section 3.13
 - lecture will cover both

4190.308 Computer Architecture, Fall 2015

15

CSE컴퓨터공학부

x86 Machine-Level Programming Basics

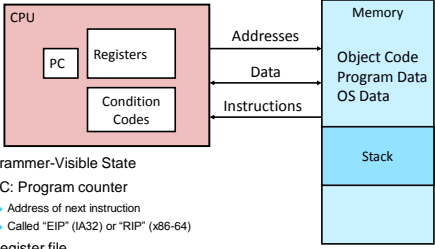
- History of Intel processors and architectures
- C, assembly, machine code
- Assembly Basics: Registers, operands, move, push/pop
- Introduction to x86-64

4190.308 Computer Architecture, Fall 2015

16

CSE컴퓨터공학부

Assembly Programmer's View



- Programmer-Visible State
 - PC: Program counter
 - Address of next instruction
 - Called "EIP" (IA32) or "RIP" (x86-64)
 - Register file
 - Heavily used program data
 - Condition codes
 - Store status information about most recent arithmetic operation
 - Used for conditional branching
- Memory
 - Byte addressable array
 - Code, user data, (some) OS data
 - Includes stack used to support procedures

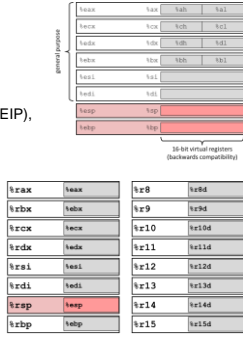
4190.308 Computer Architecture, Fall 2015

17

CSE컴퓨터공학부

x86 ISA

- x86 (32 bit):
 - register-memory architecture
 - 2 operands, max 1 memory
 - registers:
 - 6(8) general purpose registers, PC (EIP), condition codes, processor status
 - byte-addressed memory, little-endian
 - 1, 2, 4, 8 byte data types
- x86_64 (64 bit):
 - same as above except
 - 16 general purpose registers
 - 16-byte data type



4190.308 Computer Architecture, Fall 2015

18

CSE컴퓨터공학부

x86 ISA

- instruction format
 - operand size specifier
 - postfix to the operation
 - b (1 byte), w (2 bytes), l (4 bytes), q (8 bytes – x86_64 only)
 - two-operand format: first source is implicit destination
 - $a = a + b$
 - in x86 assembly:

addl b, a
 - careful with non-commutative operations:

subl a, b ≠ subl b, a

4190.308 Computer Architecture, Fall 2015

19

CSE 컴퓨터공학부

x86 ISA

- Operand specifiers

Type	Form	Operand value	Name
Immediate	$\$Imm$	Imm	Immediate
Register	Ea	$R[Ea]$	Register
Memory	Imm	$M[Imm]$	Absolute
	(Eb)	$M[R(Eb)]$	Indirect
	$Imm(Eb)$	$M[R(Eb) + Imm]$	Base + displacement
	$(Eb, E1)$	$M[R(Eb) + R(E1)]$	Indexed
	$Imm(Eb, E1)$	$M[R(Eb) + R(E1) + Imm]$	Indexed
	$(, E1, s)$	$M[R(E1) * s]$	Scaled indexed
	$Imm(, E1, s)$	$M[R(E1) * s + Imm]$	Scaled indexed
	$(Eb, E1, s)$	$M[R(Eb) + R(E1) * s]$	Scaled indexed
	$Imm(Eb, E1, s)$	$M[R(Eb) + R(E1) * s + Imm]$	Scaled indexed

4190.308 Computer Architecture, Fall 2015

20

CSE 컴퓨터공학부

From C to Machine Code

- Code in files `p1.c p2.c`
- Compile with command: `gcc -O p1.c p2.c -o p`
 - Use optimizations (`-O`)
 - Put resulting binary in file `p`

text

C program (p1.c p2.c)

↓ Compiler (gcc -S)

text

asm program (p1.s p2.s)

↓ Assembler (gcc or as)

binary

object program (p1.o p2.o)

↓ Linker (gcc or ld)

static libraries (.a)

↓

binary

executable program (p)

4190.308 Computer Architecture, Fall 2015

21

CSE 컴퓨터공학부

Compiling To IA32 Assembly

C Code

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```

`$ gcc -m32 -O -S code.c`
produces the file `code.s`:

Generated IA32 Assembly by gcc-4.4.6 on CentOS 6.3

```
sum:
    pushl   %ebp
    movl    %esp,%ebp
    movl    12(%ebp),%eax
    addl    8(%ebp),%eax
    popl    %ebp
    ret
```

→ compare output generated with `gcc -m32 -O0 -S code.c`

4190.308 Computer Architecture, Fall 2015

22

CSE 컴퓨터공학부

Compiling To IA32 Assembly – GCC Version

C Code

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```

`$ gcc -m32 -O -S code.c`
produces the file `code.s`:

Generated IA32 Assembly by gcc-4.6.3 on Gentoo

```
sum:
.LFB0:
.cfi_startproc
movl    8(%esp), %eax
addl    4(%esp), %eax
ret
.cfi_endproc
```

→ compare output generated with `gcc -m32 -O0 -S code.c`

4190.308 Computer Architecture, Fall 2015

23

CSE 컴퓨터공학부

Compiling To x86_64 Assembly

C Code

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```

`$ gcc -m64 -O -S code.c`
produces the file `code.s`:

Generated x86_64 assembly (identical for gcc-4.4.6/gcc-4.6.3)

```
sum:
.LFB0:
.cfi_startproc
leal    (%rsi,%rdi), %eax
ret
.cfi_endproc
```

→ compare output generated with `gcc -m64 -O0 -S code.c`

4190.308 Computer Architecture, Fall 2015

24

CSE 컴퓨터공학부

Assembly Characteristics: Data Types

- "Integer" data of 1, 2, or 4 bytes
 - Data values
 - Addresses (untyped pointers)
- Floating point data of 4, 8, or 10 bytes
- No aggregate types such as arrays or structures
 - Just contiguously allocated bytes in memory

4190.308 Computer Architecture, Fall 2015

25



Assembly Characteristics: Operations

- **Arithmetic/Logic Operations**
Perform arithmetic/logic functions on registers or memory data
- **Memory Operations**
Transfer data between memory and registers
 - Load data from memory into register
 - Store register data into memory
- **Control Transfer Operations**
Transfer control
 - Unconditional jumps to/from procedures
 - Conditional branches
- **Special Operations**
 - processor control, loop operations, ...

4190.308 Computer Architecture, Fall 2015

26



Object Code

IA32 Code for sum

```
0x401040 <sum>:
0x55
0x89
0xe5
0x8b
0x45
0x0c
0x03
0x45
0x08
0x5d
0xc3
```

- Total of 11 bytes
- Each instruction 1, 2, or 3 bytes
- Starts at address 0x401040

- **Assembler**
 - Translates .s into .o
 - Binary encoding of each instruction
 - Nearly-complete image of executable code
 - Missing linkages between code in different files
- **Linker**
 - Resolves references between files
 - Combines with static run-time libraries
 - E.g., code for `malloc`, `printf`
 - Some libraries are *dynamically linked*
 - Linking occurs when program begins execution

4190.308 Computer Architecture, Fall 2015

27



Machine Instruction Example

```
int t = x+y;
```

- C Code
 - Add two signed integers

```
addl 8(%ebp), %eax
```

Similar to expression:

```
x += y
```

More precisely:

```
int eax;
int *ebp;
eax += ebp[2]
```

- Assembly
 - Add two 4-byte integers
 - "long" words in GCC parlance
 - same instruction whether signed or unsigned
 - Operands:
 - x:** Register `%eax`
 - y:** Memory `M[%ebp+8]`
 - t:** Register `%eax`
 - return function value in `%eax`

```
0x0804809a: 03 45 08
```

- Object Code
 - 3-byte instruction
 - Stored at address `0x0804809a`

4190.308 Computer Architecture, Fall 2015

28



Disassembling Object Code

```
$ gcc -m32 -c -O code.c
$ objdump -d code.o
```

```
00000000 <sum>:
0: 55      push    %ebp
1: 89 e5   mov     %esp, %ebp
3: 8b 45 0c mov     0xc(%ebp), %eax
6: 03 45 08 add     0x8(%ebp), %eax
9: 5d      pop     %ebp
a: c3      ret
```

- Disassembler
 - Useful tool for examining object code
 - Analyzes bit pattern of series of instructions
 - Produces *approximate* rendition of assembly code
 - Can be run on either `a.out` (complete executable) or `.o` file

4190.308 Computer Architecture, Fall 2015

29



Alternate Disassembly using gdb

- Within `gdb` Debugger

```
$ gcc -m32 -O -o p code.c
$ gdb p
```

```
(gdb) disassemble sum
Dump of assembler code for function sum:
0x08048394 <+0>: push    %ebp
0x08048395 <+1>: mov     %esp, %ebp
0x08048397 <+3>: mov     0xc(%ebp), %eax
0x0804839a <+6>: add     0x8(%ebp), %eax
0x0804839d <+9>: pop     %ebp
0x0804839e <+10>: ret
End of assembler dump.
(gdb) x/11xb sum
0x08048394 <sum>: 0x55 0x89 0xe5 0x8b 0x45 ...
0x0804839c <sum+8>: 0x08 0x5d 0xc3
(gdb)
```

4190.308 Computer Architecture, Fall 2015

30



What Can be Disassembled?

```
% objdump -d WINWORD.EXE

WINWORD.EXE:  file format pei-i386

No symbols in "WINWORD.EXE".
Disassembly of section .text:

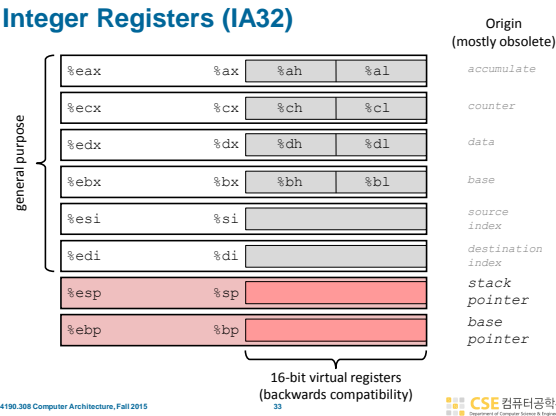
30001000 <.text>:
30001000:  55          push  %ebp
30001001:  8b ec       mov   %esp,%ebp
30001003:  6a ff       push  $0xffffffff
30001005:  68 90 10 00 push  $0x30001090
3000100a:  68 91 dc 4c 30 push  $0x304cdc91
```

- Anything that can be interpreted as executable code (not everything makes sense, of course)
- Disassembler examines bytes and reconstructs assembly source

x86 Machine-Level Programming Basics

- History of Intel processors and architectures
- C, assembly, machine code
- **Assembly Basics: Registers, operands, move, push/pop**
- Introduction to x86-64

Integer Registers (IA32)



Moving Data: IA32

- Moving Data
- **movx Source, Dest**
- with x in {b, w, l}
- movl source, dest
- move 4-byte "long word"
- movw source, dest
- move 2-byte "word"
- movb source, dest
- move 1-byte "byte"
- lots of move instructions in typical code

Moving Data: IA32

- Moving Data
- **movl Source, Dest**
- Operand Types
- **Immediate:** Constant integer data
- Example: `$0x400, $-533`
- Like C constant, but prefixed with `'$'`
- Encoded with 1, 2, or 4 bytes
- **Register:** One of 8 integer registers
- Example: `%eax, %edx`
- But `%esp` and `%ebp` reserved for special use
- Others have special uses for particular instructions
- **Memory:** 4 consecutive bytes of memory at address given by register
- Simplest example: `(%eax)`
- Various other "address modes"

movl Operand Combinations

	Source	Dest	Src, Dest	C Analog
movl	Imm	Reg	movl \$0x4, %eax	temp = 0x4;
		Mem	movl \$-147, (%eax)	*p = -147;
	Reg	Reg	movl %eax, %edx	temp2 = temp1;
		Mem	movl %eax, (%edx)	*p = temp;
	Mem	Reg	movl (%eax), %edx	temp = *p;
		Mem		

Cannot do memory-memory transfer with a single instruction

Simple Memory Addressing Modes

■ Normal

(R)

Mem[Reg[R]]

• Register R specifies memory address

movl (%ecx), %eax

■ Displacement

D(R)

Mem[Reg[R]+D]

• Register R specifies start of memory region

• Constant displacement D specifies offset

movl 8(%ebp), %edx

4190.308 Computer Architecture, Fall 2015

37

CSE 컴퓨터공학부

Using Simple Addressing Modes

void swap(int *xp, int *yp)

{
 int t0 = *xp;
 int t1 = *yp;
 *xp = t1;
 *yp = t0;
}

swap:

pushl %ebp
movl %esp, %ebp
pushl %ebx

movl 8(%ebp), %edx
movl 12(%ebp), %ecx
movl (%edx), %ebx
movl (%ecx), %eax
movl %eax, (%edx)
movl %ebx, (%ecx)

popl %ebx
popl %ebp
ret

} Set Up

} Body

} Finish

4190.308 Computer Architecture, Fall 2015

38

CSE 컴퓨터공학부

Understanding Swap

void swap(int *xp, int *yp)

{
 int t0 = *xp;
 int t1 = *yp;
 *xp = t1;
 *yp = t0;
}

Register	Value
%edx	xp
%ecx	yp
%ebx	t0
%eax	t1

Stack (in memory)

Offset

12 yp

8 xp

4 return adr

0 old %ebp

-4 old %ebx

%ebp

%esp

movl 8(%ebp), %edx # edx = xp

movl 12(%ebp), %ecx # ecx = yp

movl (%edx), %ebx # ebx = *xp (t0)

movl (%ecx), %eax # eax = *yp (t1)

movl %eax, (%edx) # *xp = t1

movl %ebx, (%ecx) # *yp = t0

4190.308 Computer Architecture, Fall 2015

39

CSE 컴퓨터공학부

Understanding Swap

Address

123 0x124

456 0x120

0x11c

0x118

0x114

0x110

0x10c

0x108

0x104

0x100

Offset

12 yp

8 xp

4 return adr

0 %ebp

-4 %esp

%eax

%edx

%ecx

%ebx

%esi

%edi

%esp

%ebp 0x104

movl 8(%ebp), %edx # edx = xp

movl 12(%ebp), %ecx # ecx = yp

movl (%edx), %ebx # ebx = *xp (t0)

movl (%ecx), %eax # eax = *yp (t1)

movl %eax, (%edx) # *xp = t1

movl %ebx, (%ecx) # *yp = t0

4190.308 Computer Architecture, Fall 2015

40

CSE 컴퓨터공학부

Understanding Swap

%eax

%edx 0x124

%ecx

%ebx

%esi

%edi

%esp

%ebp 0x104

Address

123 0x124

456 0x120

0x11c

0x118

0x114

0x110

0x10c

0x108

0x104

0x100

Offset

12 yp

8 xp

4 return adr

0 %ebp

-4 %esp

movl 8(%ebp), %edx # edx = xp

movl 12(%ebp), %ecx # ecx = yp

movl (%edx), %ebx # ebx = *xp (t0)

movl (%ecx), %eax # eax = *yp (t1)

movl %eax, (%edx) # *xp = t1

movl %ebx, (%ecx) # *yp = t0

4190.308 Computer Architecture, Fall 2015

41

CSE 컴퓨터공학부

Understanding Swap

%eax

%edx 0x124

%ecx 0x120

%ebx

%esi

%edi

%esp

%ebp 0x104

Address

123 0x124

456 0x120

0x11c

0x118

0x114

0x110

0x10c

0x108

0x104

0x100

Offset

12 yp

8 xp

4 return adr

0 %ebp

-4 %esp

movl 8(%ebp), %edx # edx = xp

movl 12(%ebp), %ecx # ecx = yp

movl (%edx), %ebx # ebx = *xp (t0)

movl (%ecx), %eax # eax = *yp (t1)

movl %eax, (%edx) # *xp = t1

movl %ebx, (%ecx) # *yp = t0

4190.308 Computer Architecture, Fall 2015

42

CSE 컴퓨터공학부

Understanding Swap

%eax	
%edx	0x124
%ecx	0x120
%ebx	123
%esi	
%edi	
%esp	
%ebp	0x104

Offset

yp 12

xp 8

4

%ebp → 0

-4

123	Address 0x124
456	0x120
	0x11c
	0x118
	0x114
0x120	0x110
0x124	0x10c
return adr	0x108
	0x104
	0x100

movl 8(%ebp), %edx # edx = xp

movl 12(%ebp), %ecx # ecx = yp

movl (%edx), %ebx # ebx = *xp (t0)

movl (%ecx), %eax # eax = *yp (t1)

movl %eax, (%edx) # *xp = t1

movl %ebx, (%ecx) # *yp = t0

4190.308 Computer Architecture, Fall 201543CSE 컴퓨터공학부

Understanding Swap

%eax	456
%edx	0x124
%ecx	0x120
%ebx	123
%esi	
%edi	
%esp	
%ebp	0x104

Offset

yp 12

xp 8

4

%ebp → 0

-4

123	Address 0x124
456	0x120
	0x11c
	0x118
	0x114
0x120	0x110
0x124	0x10c
return adr	0x108
	0x104
	0x100

movl 8(%ebp), %edx # edx = xp

movl 12(%ebp), %ecx # ecx = yp

movl (%edx), %ebx # ebx = *xp (t0)

movl (%ecx), %eax # eax = *yp (t1)

movl %eax, (%edx) # *xp = t1

movl %ebx, (%ecx) # *yp = t0

4190.308 Computer Architecture, Fall 201544CSE 컴퓨터공학부

Understanding Swap

%eax	456
%edx	0x124
%ecx	0x120
%ebx	123
%esi	
%edi	
%esp	
%ebp	0x104

Offset

yp 12

xp 8

4

%ebp → 0

-4

456	Address 0x124
456	0x120
	0x11c
	0x118
	0x114
0x120	0x110
0x124	0x10c
return adr	0x108
	0x104
	0x100

movl 8(%ebp), %edx # edx = xp

movl 12(%ebp), %ecx # ecx = yp

movl (%edx), %ebx # ebx = *xp (t0)

movl (%ecx), %eax # eax = *yp (t1)

movl %eax, (%edx) # *xp = t1

movl %ebx, (%ecx) # *yp = t0

4190.308 Computer Architecture, Fall 201545CSE 컴퓨터공학부

Understanding Swap

%eax	456
%edx	0x124
%ecx	0x120
%ebx	123
%esi	
%edi	
%esp	
%ebp	0x104

Offset

yp 12

xp 8

4

%ebp → 0

-4

456	Address 0x124
123	0x120
	0x11c
	0x118
	0x114
0x120	0x110
0x124	0x10c
return adr	0x108
	0x104
	0x100

movl 8(%ebp), %edx # edx = xp

movl 12(%ebp), %ecx # ecx = yp

movl (%edx), %ebx # ebx = *xp (t0)

movl (%ecx), %eax # eax = *yp (t1)

movl %eax, (%edx) # *xp = t1

movl %ebx, (%ecx) # *yp = t0

4190.308 Computer Architecture, Fall 201546CSE 컴퓨터공학부

Complete Memory Addressing Modes

■ General Form:

D (Rb, Ri, S)

Mem[Reg[Rb] + S*Reg[Ri] + D]

• D: Constant "displacement" 1, 2, or 4 bytes

• Rb: Base register: Any of 8 integer registers

• Ri: Index register: Any, except for %esp

- ↳ Unlikely you'd use %ebp, either

• S: Scale: 1, 2, 4, or 8

Discussed later in detail

4190.308 Computer Architecture, Fall 201547CSE 컴퓨터공학부

Sign/Zero Expansion

■ MOV with operands of different sizes

void main(void){char c = -1;int i = c;printf("%d\n", i);}

VS

void main(void){unsigned char c = -1;int i = c;printf("%d\n", i);}

• expected result?

4190.308 Computer Architecture, Fall 201548CSE 컴퓨터공학부

CSE 컴퓨터공학부

Department of Computer Science & Engineering

8

Sign/Zero Expansion

■ **MOV** with operands of different sizes

```
void main(void)
{
    char c = -1;
    int i = c;
    printf("%d\n", i);
}
```

VS

```
void main(void)
{
    unsigned char c = -1;
    int i = c;
    printf("%d\n", i);
}
```

• **disassembly**

```
...      movb    $0xff, -0x5(%rbp)
movzbl  -0x5(%rbp), %eax
...
```

```
...      movb    $0xff, -0x5(%rbp)
movzbl  -0x5(%rbp), %eax
...
```


Sign/Zero Expansion

■ **MOV** with operands of different sizes

```
void main(void)
{
    char c = -1;
    int i = c;
    printf("%d\n", i);
}
```

VS

```
void main(void)
{
    unsigned char c = -1;
    int i = c;
    printf("%d\n", i);
}
```



Sign/Zero Expansion

■ **MOV** with operands of different sizes

- **mov**z**bl**: move **z**ero-expand **b**yte to double word(**l**)
- **mov**s**bl**: move **s**ign-expand **b**yte to double word(**l**)

Instruction	S, D	Effect	Description
MOV	S, D	D ← S	Move (equal-sized operands)
movb		move byte	
movw		move word (16-bit)	
movl		move double word (32-bit)	
movq		move quad word (64-bit)	
MOVS	S, D	D ← SignExtend(S)	Move with sign extension
movsb[w,l,q]		move sign-extended byte to word, double word, quad word	
movsw[l,q]		move sign-extended word to double word, quad word	
movslq		move sign-extended double word to quad word	
MOVZ	S, D	D ← ZeroExtend(S)	Move with zero extension
movzb[w,l,q]		move zero-extended byte to word, double word, quad word	
movzw[l,q]		move zero-extended word to double word, quad word	
movzslq		move zero-extended double word to quad word	

- **movs/z**: always size(source operand) < size(destination operand)

Stack Operations: PUSH & POP

- **%esp/%rsp** point to a memory region call the **stack**
- the stack holds temporary data associated with currently active functions
- **push/pop** are used to store/retrieve temporary data onto/from the stack

Instruction	Effect	Description
STACK		Stack operations
pushl S	$R[\%esp] \leftarrow R[\%esp] - 4;$ $M[R[\%esp]] \leftarrow S$	push double word onto stack
popl D	$D \leftarrow M[R[\%esp]];$ $R[\%esp] \leftarrow R[\%esp] + 4$	pop double word from stack

- **pushl **s**** has same effect as "**subl \$4, %esp; movl **s**, (%esp)**"
- **popl **d**** has same effect as "**movl (%esp), **d**; addl \$4, %esp**"

- size modifiers apply: **push[b,w,l,q], pop[b,w,l,q]**

Stack Operations: PUSH & POP

- **%esp/%rsp** point to a memory region call the **stack**
- the stack holds temporary data associated with currently active functions
- **push/pop** are used to store/retrieve temporary data onto/from the stack

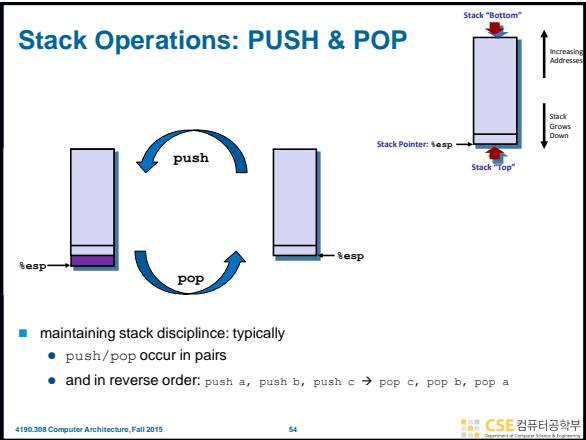
Instruction	Effect	Description
STACK		Stack operations
pushl S	$R[\%esp] \leftarrow R[\%esp] - 4;$ $M[R[\%esp]] \leftarrow S$	push double word onto stack
popl D	$D \leftarrow M[R[\%esp]];$ $R[\%esp] \leftarrow R[\%esp] + 4$	pop double word from stack

- **pushl **s**** has same effect as "**subl \$4, %esp; movl **s**, (%esp)**"
- **popl **d**** has same effect as "**movl (%esp), **d**; addl \$4, %esp**"

- size modifiers apply: **push[b,w,l,q], pop[b,w,l,q]**

Stack Operations: PUSH & POP

- maintaining stack discipline: typically
 - **push/pop** occur in pairs
 - and in reverse order: **push a, push b, push c → pop c, pop b, pop a**



x86 Machine-Level Programming Basics

- History of Intel processors and architectures
- C, assembly, machine code
- Assembly Basics: Registers, operands, move, push/pop
- Introduction to x86-64

Data Representations: IA32 + x86-64

■ Sizes of C Objects (in Bytes)

C Data Type	Generic 32-bit	Intel IA32	x86-64
unsigned	4	4	4
int	4	4	4
long int	4	4	8
char	1	1	1
short	2	2	2
float	4	4	4
double	8	8	8
long double	8	10/12	16
char *	4	4	8

– Or any other pointer


x86-64 Integer Registers

%rax	%eax	%r8	%r8d
%rbx	%ebx	%r9	%r9d
%rcx	%ecx	%r10	%r10d
%rdx	%edx	%r11	%r11d
%rsi	%esi	%r12	%r12d
%rdi	%edi	%r13	%r13d
%rsp	%esp	%r14	%r14d
%rbp	%ebp	%r15	%r15d

- Extend existing registers. Add 8 new ones.
- Make %ebp/%rbp general purpose

Instructions

- Long word l (4 Bytes) ↔ Quad word q (8 Bytes)
- New operand specifier: q
 - movq
 - addq
 - salq
 - etc.
- Note: 32-bit instructions that generate 32-bit results
 - Set higher order bits of destination register to 0
 - Example: addl \$-1, %eax



32-bit code for swap

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

```
swap:
    pushl %ebp
    movl %esp, %ebp
    pushl %ebx

    movl 8(%ebp), %edx
    movl 12(%ebp), %ecx
    movl (%edx), %ebx
    movl (%ecx), %eax
    movl %eax, (%edx)
    movl %ebx, (%ecx)

    popl %ebx
    popl %ebp
    ret
```

Set Up

Body

Finish

64-bit code for swap

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

```
swap:
    movl (%rdi), %edx
    movl (%rsi), %eax
    movl %eax, (%rdi)
    movl %edx, (%rsi)

    ret
```

Set Up

Body

Finish

- Operands passed in registers (why is that useful?)
 - First (xp) in %rdi, second (yp) in %rsi
 - 64-bit pointers
- No stack operations required
- 32-bit data
 - Data held in registers %eax and %edx
 - movl operation

64-bit code for long int swap

```
void swap(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

```
swap_1:
    movq    (%rdi), %rdx
    movq    (%rsi), %rax
    movq    %rax, (%rdi)
    movq    %rdx, (%rsi)
    ret
```

Set Up


Body

Finish

- 64-bit data
 - Data held in registers `%rax` and `%rdx`
 - `movq` operation
 - "q" stands for quad-word

Machine-Level Programming Basics: Summary

- History of Intel processors and architectures
 - Evolutionary design leads to many quirks and artifacts
- C, assembly, machine code
 - Compiler must transform statements, expressions, procedures into low-level instruction sequences
- Assembly Basics: Registers, operands, move
 - The x86 move instructions cover wide range of data movement forms
- Intro to x86-64
 - A major departure from the style of code seen in IA32

CSE 컴퓨터공학부
Department of Computer Science & Engineering

11